

Education

- **Mila, University of Montreal** Montreal, Canada
PhD in Reinforcement Learning, supervised by Aaron Courville *Fall 2017 - May 2022*
- **Indian Institute of Technology (IIT), Kharagpur** Kharagpur, India
Integrated M.Sc. in Mathematics and Computing *2011 - 2016*

Work Experience

- **Google DeepMind** London, UK
Research Scientist, Senior Research Scientist *October 2022 - Present*
 - Core Contributor on Gemini 1.0, Gemini 1.5 and Gemini 2.0; worked across several verticals: Evals, Post-Training, Reasoning, Computer-Use.
 - Co-led key RL (Reinforcement Learning) / Inference-time breakthroughs that led to Gemini's largest improvements till date across reasoning and coding benchmarks. These were shipped as part of our flagship models, **Gemini 2.0 Flash** and **Gemini 2.0 Flash Thinking**.
 - Contributed a series of post-training quality improvements that led to **Gemini 1.5 Pro** being #1 on LMSYS.
 - Made key RL contributions to Gemini's computer use agent **Project Mariner**: a frontier model/agent with state-of-the-art performance on web navigation tasks.
 - Built internal, leaked evals to track frontier reasoning and long context capabilities.
- **DeepMind** London, UK
Research Intern, Advisor: Jessica Hamrick *Feb 2021-June 2021*
 - Worked on identifying generalization capabilities of MuZero, and how to amplify them using self-supervised learning.
 - The internship project led to a first-author paper at ICLR'22: **Procedural Generalization by Planning with Self-Supervised World Models**
- **Microsoft Research** Montreal, Canada
Research Intern, Advisor: Philip Bachman *Sept 2019-May 2020*
 - Worked on self-supervised learning for data-efficient RL.
 - The internship project led to a first-author publication at ICLR'21: **Data-Efficient Reinforcement Learning with Self-Predictive Representations**
- **Microsoft Research** Montreal, Canada
Research Intern, Advisor: Marc Cote, Devon Hjelm *January-June 2019*
 - Research at the intersection of Reinforcement Learning and Representation Learning.
 - The internship project led to a first-author publication at NeurIPS'19: **Unsupervised State Representation Learning in Atari**
- **VISA Inc.** Bangalore, India
Software Engineer *August 2016-August 2017*
 - Full stack development for the VISA Developer Platform

- **Google Summer of Code**

- *Student Developer*

Remote

May-August 2015

- Built an online analytics platform for BRL-CAD which provides aggregated analytics for logs and performance metrics collected across different machines and platforms.

Selected Publications

- **Beyond human data: Scaling self-training for problem-solving with LLMs** ([arxiv](#))
*Avi Singh, John D Co-Reyes, Rishabh Agarwal, **Ankesh Anand**, Piyush Patil, Xavier Garcia, ..., Ethan Dyer, Behnam Neyshabur, Jascha Sohl-Dickstein, Noah Fiedel*
TMLR
- **Procedural Generalization by Planning with Self-Supervised World Models** ([arxiv](#))
Ankesh Anand, Jacob Walker*, Yazhe Li, Eszter Vrtes, Julian Schrittwieser, Sherjil Ozair, Thophane Weber, Jessica B. Hamrick*
Internal Conference on Learning Representations (ICLR), 2022
- **Pretraining Representations for Data-Efficient Reinforcement Learning** ([arxiv](#))
*Max Schwarzer, Nitarshan Rajkumar, Michael Noukhovitch, **Ankesh Anand**, Laurent Charlin, Devon Hjelm, Philip Bachman, Aaron Courville*
Neural Information Processing Systems (NeurIPS), 2021
- **Data-Efficient Reinforcement Learning with Self-Predictive Representations** ([arxiv](#))
Max Schwarzer, **Ankesh Anand***, Rishab Goel, R Devon Hjelm, Aaron Courville, Philip Bachman*
Internal Conference on Learning Representations (ICLR), 2021 (Spotlight)
- **Unsupervised State Representation Learning in Atari** ([arxiv](#))
Ankesh Anand, Evan Racah*, Sherjil Ozair*, Yoshua Bengio, Marc Cote, Devon Hjelm*
Neural Information Processing Systems (NeurIPS), 2019
- **HoME: a Household Multimodal Environment** ([arxiv](#))
Simon Brodeur, Ethan Perez, **Ankesh Anand***, Florian Golemo*, Luca Celotti, Florian Strub, Jean Rouat, Hugo Larochelle, Aaron Courville*
International Conference on Learning Representations (ICLR) Workshop Track, 2018
- **We used Neural Networks to Detect Clickbaits: You won't believe what happened Next!** ([arxiv](#))
***Ankesh Anand**, Tanmoy Chakraborty, Noseong Park*
European Conference on Information Retrieval (ECIR), 2017

Technical Skills

- **Programming Languages:**
 - **Proficient:** Python, JavaScript, C++
- **Machine Learning Libraries:** PyTorch, JAX, numpy, Pandas
- **Web Development:** Django, Flask, NodeJS, ReactJS, HTML5, CSS3, MySQL